

Use custom address for listening socket

#15988

New issue

Merged

njhill merged 6 commits into vllm-project:main from jglaser:use_socket_with_addr on Apr 25, 2025

Conversation7

Commits6

Checks4

Files changed1

+21-1



jglaser commented on Apr 3, 2025

Contributor

• edited by github-actions (bot)

This sets up TCPStore to listen on a specified interface instead of on 0.0.0.0, which saves OS resources

Currently, using the old TCP store (`use_libuv=False`), as the new one (`libuv=True`) only supports custom socket setup in the most recent versions of Pytorch.

FIX [#15987](#) (link existing issues this PR will resolve)

Reviewers

- youkaichao
- njhill

Assignees

No one assigned

Labels

ready

Projects

None yet

Milestone

No milestone

Development

Successfully merging this pull request may close these issues.

[Bug]: EADDRINUSE (-98) error when s...

3 participants



github-actions (bot) commented on Apr 3, 2025

Hi! Thank you for contributing to the vLLM project.

Join our developer Slack at <https://slack.vllm.ai> to discuss your PR in #pr-reviews, coordinate on features in #feat-channels, or join special interest groups in #sig- channels.

Just a reminder: PRs would not trigger full CI run by default. Instead, it would only run `fastcheck` CI which starts running only a small and essential subset of CI tests to quickly catch errors. You can run other CI tests on top of those by going to your `fastcheck` build on Buildkite UI (linked in the PR checks section) and unblock them. If you do not have permission to unblock, ping `simon-mo` or `khluu` to add you in our Buildkite org.

Once the PR is approved and ready to go, your PR reviewer(s) can run CI to test the changes comprehensively before merging.

To run CI, PR reviewers can either: Add `ready` label to the PR or enable auto-merge.



mgoin requested review from **njhill** and **youkaichao**
10 months ago



njhill approved these changes
on Apr 3, 2025

[View reviewed changes](#)

njhill left a comment

Member

Thanks [@jglaser](#), looks great.

Could you fix the linter error? (comment line too long)



jglaser commented on Apr 5, 2025

Contributor

Author

[@njhill](#) I am not sure why the two remaining tests are failing
... might be unrelated ... mind having a look?



WoosukKwon requested a review from **njhill**
10 months ago



WoosukKwon commented on Apr 5, 2025

Collaborator

[@njhill](#) Does this PR need force merge?



njhill commented on Apr 8, 2025

Member

[@jglaser](#) could you rebase on latest main branch and add a
sign-off to your commits for the DCO check ☐ :

[https://github.com/vllm-project/vllm/pull/15988/checks?
check_run_id=39931987909](https://github.com/vllm-project/vllm/pull/15988/checks?check_run_id=39931987909)



1



jglaser added 5 commits 9 months ago



Use custom address for listening
socket ...

d70a354



Fix typing error ...

08fe582

-   [Fix too long line](#) ... [5d024b1](#)
-   [fix another long comment](#) ... [73ad6ca](#)
-   [fix remaining pre-commit errors](#) ...  [eb02f37](#)

  **jglaser** force-pushed the [use_socket_with_addr](#) branch from [3d348ab](#) to [eb02f37](#) [9 months ago](#) [Compare](#)

  **njhill** added the [ready](#) label [on Apr 16, 2025](#)

  **njhill** enabled auto-merge (squash) [9 months ago](#)



njhill commented [on Apr 17, 2025](#)

Member

[@jglaser](#) would you mind merging in latest main (hopefully) one last time?

 1

  Merge branch 'main' into [use_socket_with_addr](#)  [9bce2a3](#)



jglaser commented [on Apr 25, 2025](#)

Contributor

Author

[@jglaser](#) would you mind merging in latest main (hopefully) one last time?

just did

 1



 **njhill** merged commit [0d6e187](#) into [vllm-project:main](#) [on Apr 25, 2025](#)

[View details](#)

49 checks passed

 **gshttras** added a commit to ROCm/vllm that referenced this pull request [on Apr 26, 2025](#)



[Upstream merge 2025 04 25](#) [\(#524\)](#) ...

Verified

 [28007b0](#)



jikunshang pushed a commit to jikunshang/vllm that referenced this pull request [on Apr 29, 2025](#)

 [Use custom address for listening socket \(vllm-project#15988\)](#) [8c4163c](#) [...](#)

 **lk-chen** pushed a commit to lk-chen/vllm that referenced this pull request [on Apr 30, 2025](#)

 [Use custom address for listening socket \(vllm-project#15988\)](#) [612d065](#) [...](#)

 **adobrzyn** pushed a commit to HabanaAI/vllm-fork that referenced this pull request [on Apr 30, 2025](#)

 [Use custom address for listening socket \(vllm-project#15988\)](#) [0989bca](#) [...](#)

 **russellb** mentioned this pull request [on May 2, 2025](#)

[\[Security\] Document StatelessProcessGroup security concerns #17591](#)

 Closed

 **RichardoMrMu** pushed a commit to RichardoMrMu/vllm that referenced this pull request [on May 12, 2025](#)

 [Use custom address for listening socket \(vllm-project#15988\)](#) [8fb258e](#) [...](#)

  **ckhordiasma** mentioned this pull request [on May 15, 2025](#)

[nm vllm ent 0.8.5 sync red-hat-data-services/vllm#139](#)

 Merged

 **russellb** added a commit to russellb/vllm that referenced this pull request [on Jun 15, 2025](#)

 [\[Security\] Document StatelessProcessGroup security concerns](#) [595e1d7](#) [...](#)

 **russellb** added a commit to russellb/vllm that referenced this pull request [on Aug 6, 2025](#)

 [\[Security\] Document StatelessProcessGroup security concerns](#) [90a22fa](#) [...](#)



dsocek pushed a commit to HabanaAI/vllm-fork that referenced this pull request on Sep 20, 2025



Use custom address for listening
socket_([vllm-project#15988](#)) ...

9e6754e