

[Core] Upgrade to xgrammar 0.1.18, add cache size limit #16283

New issue

Merged

vllm-bot merged 1 commit into `vllm-project:main` from `russellb:xgrammar-0.1.18` on Apr 9, 2025

Conversation2

Commits1

Checks7

Files changed4

+22-3



russellb commented on Apr 9, 2025

Member

xgrammar 0.1.18 includes a fix for security vulnerability where a malicious user can cause vllm to consume all of the host's RAM with the knowledge that each unique grammar adds an entry to an unbounded cache in memory.

The cache was implemented in xgrammar PR 243: [mlc-ai/xgrammar#243](#)

This commit makes use of the new option to limit the cache size.

The default is 512MB, which should be reasonable for most use cases.

It will cache roughly 1000 json schemas. The cache size can be changed using the `VLLM_XGRAMMAR_CACHE_MB` environment variable.

The security vulnerability was reported here: [GHSA-389x-67px-mjg3](#)

Signed-off-by: Russell Bryant rbryant@redhat.com

4

Reviewers

mgoin

✓

Assignees

No one assigned

Labels

ci/build

ready

structured-output

v1

Projects

None yet

Milestone

No milestone

Development

Successfully merging this pull request may close these issues.

None yet

3 participants

[Core] Upgrade to xgrammar 0.1.18, add cache size limit ...

✗ 549f429

russellb requested a review from mgoin as a code owner 9 months ago



github-actions bot commented on Apr 9, 2025

Hi! Thank you for contributing to the vLLM project.

Join our developer Slack at <https://slack.vllm.ai> to discuss your PR in #pr-reviews, coordinate on features in #feat-channels, or join special interest groups in #sig- channels.

Just a reminder: PRs would not trigger full CI run by default. Instead, it would only run `fastcheck` CI which starts running only a small and essential subset of CI tests to quickly catch errors. You can run other CI tests on top of those by going to your `fastcheck` build on Buildkite UI (linked in the PR checks section) and unblock them. If you do not have permission to unblock, ping `simon-mo` or `khluu` to add you in our Buildkite org.

Once the PR is approved and ready to go, your PR reviewer(s) can run CI to test the changes comprehensively before merging.

To run CI, PR reviewers can either: Add `ready` label to the PR or enable auto-merge.



mergify bot added `ci/build` `structured-output` `v1` labels on Apr 9, 2025



mgoin approved these changes on Apr 9, 2025

[View reviewed changes](#)

mgoin left a comment

Member

LGTM, thanks for the fix



mgoin added the `ready` label on Apr 9, 2025



mgoin enabled auto-merge (squash) 9 months ago



vllm-bot merged commit **cb84e45** into `vllm-project:main` on Apr 9, 2025
83 of 86 checks passed

[View details](#)



yma11 pushed a commit to yma11/vllm that referenced this pull request on Apr 16, 2025

 [\[Core\] Upgrade to xgrammar 0.1.18, add cache size limit \(vllm-project...\)](#) [bd9db6a](#)

...

 **yangw-dev** pushed a commit to yangw-dev/vllm that referenced this pull request on Apr 22, 2025

 [\[Core\] Upgrade to xgrammar 0.1.18, add cache size limit \(vllm-project...\)](#) [fe98dff](#)

...

 **lk-chen** mentioned this pull request on Apr 28, 2025

Bump xgrammar from 0.1.16 to 0.1.18 in /python ray-project/ray#52175

 Closed

 **jikunshang** pushed a commit to jikunshang/vllm that referenced this pull request on Apr 29, 2025

 [\[Core\] Upgrade to xgrammar 0.1.18, add cache size limit \(vllm-project...\)](#) [94efac2](#)

...

 **lk-chen** pushed a commit to lk-chen/vllm that referenced this pull request on Apr 30, 2025

 [\[Core\] Upgrade to xgrammar 0.1.18, add cache size limit \(vllm-project...\)](#) [3f6fcdf](#)

...

 **RichardoMrMu** pushed a commit to RichardoMrMu/vllm that referenced this pull request on May 12, 2025

 [\[Core\] Upgrade to xgrammar 0.1.18, add cache size limit \(vllm-project...\)](#) [cb7f374](#)

...

 **yma11** pushed a commit to yma11/vllm that referenced this pull request on Jun 4, 2025

 [\[Core\] Upgrade to xgrammar 0.1.18, add cache size limit \(vllm-project...\)](#) [cba80cf](#)

...

 **Wanli-Jiang** mentioned this pull request on Jun 19, 2025

feat: TRTLLM-5941 Upgrade

xgrammar to 0.1.18 NVIDIA/TensorRT-

LLM#5364

 Merged